

Model Fine-Tuning & Quarantine Architecture

PRIVILEGED BRAIN • RUN 7 • RELEASE GOVERNANCE • TRUST BOUNDARIES

THESIS

We did not train a general assistant. We specialized a 1.5B code model into a narrow, refusal-aware command policy, then placed it behind a Controller that treats model output as untrusted until grammar, schema, risk, human-approval, verifier, sandbox, and audit checks pass.

Final artifact	Base	Adaptation	Release result
run7_cot_q4km.gguf	Qwen2.5-Coder-1.5B-Instruct	LoRA SFT, structured rationale format	100% refusal • 95.5% MCP JSON • 940 MB

WHAT WAS TRAINED

The Privileged Brain (PB): local, narrow, tool-call generation and refusal behavior.

WHAT WAS QUARANTINED

The Quarantined Brain (QB): content-understanding model with tools disabled; production config uses Gemini 2.5 Flash, not the PB fine-tune.

Prepared for deep technical interview discussion • repository-grounded • September 2026

The answer to give before any details

This page prevents the two most damaging interview mistakes: claiming we fine-tuned the QB, or claiming the model itself is the security boundary.

INTERVIEW ANSWER

"We fine-tuned the Privileged Brain, not the production Quarantined Brain. The PB is a Qwen2.5-Coder-1.5B model adapted with rank-8 LoRA SFT on a safety-enriched NL → Bash corpus. The QB is isolated by architecture: tools disabled, structured intent output, Controller validation, risk escalation, HITL, and a verifier. The final model passed 20/20 adversarial refusals and 95.5% structural MCP validity after Q4_K_M quantization."

Four claims you can defend

- The training target was deliberately narrow: one safe command, REFUSE, or CLARIFY—not open-ended helpfulness.
- The biggest improvement came from changing the target representation to a two-line structured rationale + command format, not merely adding more epochs.
- We released on a vector of metrics. Run 8 improved exact match and token-F1 but regressed safety from 100% to 90%, so we rejected it.
- Model behavior is one defense layer. Deterministic controls—schema validation, risk floors, approval, sandboxing, and append-only audit—carry the security guarantee.

Claims to avoid

Do not say	Why	Say instead
"We fine-tuned both brains."	No production QB fine-tune is evidenced.	PB was fine-tuned; QB was quarantined and prompt/schema constrained.
"95.5% means tool correctness."	The harness primarily checks JSON shape and catalogue membership.	95.5% structural validity; semantic correctness is separately schema/verifier checked.
"The CoT proves the model reasons."	Rationales were supervised/deterministically generated.	Structured rationale target improved behavior; it is not proof of latent reasoning.
"The PB sees only a UUID."	Current session code sends a richer envelope.	PB excludes raw conversation but receives structured target/content/hint fields.

EVIDENCE docs/implementation_plan.md::Phase 4 close-out; privileged-brain/v2/scripts/sft_train.py::main; dual-brain/controller/session.py::SessionState.build_pb_user_turn

Fine-tuning and quarantine solve different problems

PB — SPECIALIZED POLICY MODEL
Input: a minimized, Controller-built envelope. Output: exactly one allowed MCP call or refusal. Local artifact: run7_cot_q4km.gguf. Optimization target: safe, minimal, structurally constrained execution.

QB — SEMANTIC PLANNER
Input: user request and selected context. Output: structured Intent Object. Production config: Gemini 2.5 Flash. Optimization mechanism in this repository: prompt/schema/retry—not a checked-in production fine-tune.

Why two models?

A content-capable model needs enough context to understand ambiguous natural language, but that same context is an injection surface. An execution-capable model needs precise tool schemas and policy, but should not ingest documents or open-ended conversation. Separating those capabilities reduces what any single compromised model can do.

Property	Quarantined Brain	Privileged Brain
Primary competence	intent extraction, decomposition, summarization	safe command/tool-call synthesis
Raw user/content access	yes, by design	no raw conversation; structured envelope only
Tool connectivity	none in BrainBackend requests	no direct mcpd attachment; Controller mediates
Production backend	Gemini 2.5 Flash (configurable registry)	local llama-server + Run 7 GGUF
Failure response	retry or reject invalid Intent Object	bounded repair/retry; deny on persistent invalidity
Security role	untrusted proposer	untrusted narrow proposer—not an authorization oracle
DANGER	Do not describe “quarantine training” as a new ML algorithm. In this project, quarantine is a capability and information-flow architecture. Training made the PB safer; isolation constrains the QB.	

EVIDENCE dual-brain/controller/backends/base.py::BrainBackend.complete; dual-brain/controller/config.py; cx-distro/distro/controller.toml; dual-brain/controller/prompts/qb_gemini.txt

04 OBJECTIVE

What the model was optimized to do

The objective was not maximum benchmark helpfulness. It was a narrow execution policy with calibrated refusal.

Supervised target

REASONING: <one sentence describing the operation and its safety consequence>
 COMMAND: <one Bash command | REFUSE: reason | CLARIFY: question>

At SFT time, the model minimizes next-token cross-entropy on ChatML-formatted sequences. Conceptually, for target tokens $y_1 \dots y_T$ and context x , the objective is:

$$L_{\text{SFT}}(\theta) = - \sum_t \log p(\theta(y_t | x, y_{<t}))$$

The repository passes the full rendered ChatML string to TRL SFTTrainer and does not configure an explicit completion-only collator. That suggests the loss covers the serialized conversation unless the pinned TRL version changes the default. In an interview, flag this as an implementation detail you would verify before reproducing the run.

Threat model embedded in the labels

- Remote-code execution patterns: curl|bash, wget|sh, eval of downloaded content.
- Destructive storage and filesystem operations: rm -rf /, raw-device writes, mkfs.
- Privilege abuse: chmod 777 on protected paths, sudoers manipulation, credential exfiltration.
- Ambiguity: refuse when unsafe; ask one targeted clarification when intent is underspecified.


TRANSFERABLE

This is cost-sensitive classification disguised as sequence generation: a false-negative refusal can corrupt a host; a false-positive refusal mostly costs user friction. The release gate should weight those errors asymmetrically.

EVIDENCE privileged-brain/v2/scripts/process_datasets.py::SYSTEM_PROMPT_COT; privileged-brain/v2/scripts/phase4_eval.py::DANGEROUS_PROMPTS

From data to a checksum-pinned GGUF

Fine-tuning → release artifact



RELEASE RULE Never trade safety for a marginal utility gain.

Run 8 improved lexical metrics but fell from 20/20 to 18/20 refusals — so it was rejected.

The actual flow

Stage	Implementation	Key invariant
Prepare	process_datasets.py	invalid, destructive non-refusal, and scraped-prose examples removed
Adapt	sft_train.py --cot	base frozen; rank-8 adapters updated
Prefer (optional)	dpo_train.py --cot	SFT model used as frozen reference
Evaluate	phase4_eval.py + eval_adversarial_gguf.py	quantized artifact must be re-tested
Fuse	fuse_lora.py --adapter run7_cot	adapter merged into Qwen weights on CPU
Package	HF → GGUF → Q4_K_M	940 MB artifact; SHA-256 pinned
Serve	llama-server	grammar + post-generation schema validation
INTERVIEW ANSWER	“The adapter is not the shipping artifact. We merged it into the exact base model, converted to GGUF, quantized to Q4_K_M, and evaluated that quantized model because quantization can change borderline logits and therefore refusal behavior.”	

EVIDENCE privileged-brain/v2/scripts/fuse_lora.py::main; privileged-brain/inference/grammar/*.gbnf; models/checksums.sha256

Dataset composition and quality controls

Artifact	Rows	What it represents
clean/train.jsonl	13,389	filtered ChatML SFT training
clean/valid.jsonl	1,488	random 10% validation split
cot/train.jsonl	13,839	two-line rationale/command training
cot/valid.jsonl	1,538	CoT validation
dpo_pairs.jsonl	500	chosen/rejected preferences
curated_pb.jsonl	1,027	hand-verified Run 8 continuation set

Observed label mix in the CoT corpus

Output class	Count	Share of 15,377
REFUSE	1,209	7.86%
CLARIFY	1,004	6.53%
command / other	13,164	85.61%

The clean corpus totals 14,877 records; the CoT corpus totals 15,377, an observed +500-row delta. The current tree also contains 500 DPO pairs, but file counts alone do not prove that the extra CoT rows and the preference set are identical.

Filtering rules that matter

- Length caps: natural language ≤ 300 characters; command ≤ 500 characters.
- Reject markdown/prose, Unicode dash corruption, prompt prefixes, duplicated commands, broken pipes, and pseudo-code ellipses.
- Reject known destructive commands unless represented as a REFUSE label.
- Seeded shuffle (42), then 90/10 train/validation split; 50-sample CoT format precheck before GPU time.

REVIEW

A random row split does not guarantee semantic isolation. Command paraphrases or template siblings can leak across splits. A stronger reproduction would group by normalized command family, template origin, and danger category before splitting.

EVIDENCE `privileged-brain/v2/scripts/process_datasets.py::[is_valid,_make_reasoning,assert_cot_format]`; line counts computed from checked-in JSONL artifacts

Why the structured rationale changed the outcome

BEFORE RUN 7

The fine-tuned model often produced tutorial prose or multi-command answers. Exact-match and execution evaluation punished this, and refusal behavior plateaued around 80%.

RUN 7 FORMAT

A fixed first line made the safety consequence explicit; a fixed second line isolated the executable decision. GBNF constrained this two-line language during shell-mode inference.

What the “reasoning” field is—and is not

Training rationales were generated deterministically from the target command and a tool-to-verb map when no rationale was supplied. They are supervised visible text, not hidden chain-of-thought and not evidence that the model discovered an independent causal proof. The mechanism is better described as a structured safety rationale or scratch field.

```
_make_reasoning(nl, bash) → deterministic sentence
assistant target → REASONING: ...
                  COMMAND: REFUSE: ... | CLARIFY: ... | command
```

Why it can still help

- It creates a stable token prefix that separates deliberation-like text from the executable field.
- Danger concepts appear immediately before the decision token REFUSE, strengthening local conditional association.
- It reduces mode confusion: tutorial prose has no legal slot in the grammar.
- The parser can discard the rationale and expose only COMMAND to later controls.


DANGER

Do not let a model-authored rationale become an authorization fact. Risk level, resolved target, diff, and approval must come from schema, policy, and COW preview—not persuasive prose.

EVIDENCE `privileged-brain/v2/scripts/process_datasets.py::_make_reasoning`; `privileged-brain/inference/grammar/bash_cot.gbnf`; `docs/Archive/V1_ROADMAP.md::Run 7`

LoRA in this run: exact configuration and intuition

LoRA: freeze W, learn a low-rank update



$W' = W + (\alpha/r)BA$ $r=8$ $\alpha=16$ $\alpha/r=2.0$ dropout=0.05 ~9.23M trainable params ~0.59%

For a frozen projection $W \in \mathbb{R}^{d_{out} \times d_{in}}$, LoRA learns $A \in \mathbb{R}^{r \times d_{in}}$ and $B \in \mathbb{R}^{d_{out} \times r}$. The forward pass uses $W'x = Wx + (\alpha/r)BAx$. With $r=8$ and $\alpha=16$, the adapter branch is scaled by 2.0.

Configuration	Value	Why it matters
Target modules	q/k/v/o_proj	attention routing and representation mixing
Target modules	gate/up/down_proj	MLP feature selection and transformation
Rank / alpha	8 / 16	low-capacity update; scale $\alpha/r=2$
Dropout	0.05	regularizes adapter path
Bias	none	fewer trainable degrees of freedom
Measured trainables	9,232,384 ~0.59%	Run 4 measurement under reused configuration
DEPTH	Targeting both attention and MLP projections gives the adapter leverage over token routing and command-domain feature transformation. The tradeoff is greater capacity—and therefore more opportunity to overwrite safety behavior—than attention-only LoRA.	

EVIDENCE [privileged-brain/v2/scripts/sft_train.py::LoraConfig](#); docs/Archive/run4_sft_results.md

09 OPTIMIZATION

SFT training loop, compute, and checkpoints

Hyperparameter	Run configuration
Base weights	Qwen/Qwen2.5-Coder-1.5B-Instruct, loaded bfloat16
Epochs	3
Micro-batch	2
Gradient accumulation	8
Effective batch	16 sequences/update
Learning rate	2×10^{-4}
Schedule	cosine; 264 warmup steps
Max sequence	512 tokens
Optimizer	AdamW (torch), weight decay 0.01
Packing	disabled
Checkpoint/eval	every 500 steps; keep 3; select lowest eval loss

With 13,839 CoT training records and effective batch 16, the nominal run is about 865 optimizer updates per epoch, or ~2,595 over three epochs. The fixed 264-step warmup is therefore about 10.2% of the nominal update count. This is an estimate because distributed sampler/drop-last/version behavior can shift exact counts.

Compute paths

The script selects Apple MPS, CUDA, or CPU. Low-memory mode halves the micro-batch to 1, restores effective batch 16 through accumulation 16, shortens sequence length to 256, enables gradient checkpointing, and caps the MPS high-watermark ratio at 0.6. Normal mode disables the MPS cap.

REPRODUCIBILITY

The base model ID is pinned by name, but the training script does not pin a Hugging Face revision or base-weight hash. For a forensic rerun, record the model commit SHA, tokenizer files, library lock, CUDA/MPS versions, full trainer state, and dataset hashes.

EVIDENCE `privileged-brain/v2/scripts/sft_train.py::main; privileged-brain/v2/data/processed/cot/train.jsonl`

10 PREFERENCE OPTIMIZATION

DPO existed—but the final Run 7 release was SFT-only

Direct Preference Optimization trains the policy to score a chosen response above a rejected response while remaining anchored to a reference policy. A common form is:

$$L_{DPO} = -E \log \sigma \left(\beta \left[(\log \pi_{\theta}(y+|x) - \log \pi_{ref}(y+|x)) - (\log \pi_{\theta}(y-|x) - \log \pi_{ref}(y-|x)) \right] \right)$$

DPO setting	Value
Pairs	500; 90/10 split, seed 42
Epochs	2
Micro-batch / accumulation	1 / 8
Learning rate	5×10 ⁻⁵
β	0.1
Max length	512
Schedule	cosine; 10% warmup

The subtle but correct reference-model choice

The DPO script loads the SFT adapter twice: one trainable policy and one adapter merged into the base and frozen as π_{ref} . Using the bare base as the reference would be misleading because it assigns very low probability to the project’s REFUSE/CoT format; the preference margin would partly reward learning the format rather than the safety preference.

ACTUAL RELEASE

DPO was used in earlier experiments and remains implemented. It was not required for the final Run 7 artifact because CoT-format SFT already reached 20/20 refusals. Do not narrate Run 7 as “SFT then DPO.”

EVIDENCE privileged-brain/v2/scripts/dpo_train.py::[main,DPOConfig]; docs/implementation_plan.md::Phase 4 close-out

11 EXPERIMENT HISTORY

Failure-driven iteration across runs

Run	Change / evidence	Utility signal	Safety signal	Decision
3	NL2Bash + 616 synthetic; SFT	92.6% token acc; 27.5% FEH	curl bash regression recorded	baseline only
4	cleaner data + SFT; optional DPO	93.41% token acc; eval loss .2928	SFT/DPO 15/20 refusals	not releasable
6	targeted iterations	93.34% token acc; 3.1% historical FEH	16/20 refusals	plateau
7	structured rationale + command SFT	7.7% exact; 53.6% token-F1; 23.3% exec-FE	20/20; 95.5% MCP JSON	released
8	1,027 curated + ~3k replay; lr 5e-5, 2 epochs	12% exact; 55.9% token-F1	18/20; executes chmod 777 /etc	rejected

What changed our mind


Runs 4–6 suggested that more of the same training signal would not clear the safety ceiling. Run 7 changed the representation, producing a clean decision boundary between rationale and executable output. Run 8 then showed catastrophic forgetting / decision-boundary drift: replay and a lower learning rate were not enough to preserve two critical refusal cases.

INTERVIEW ANSWER	"Our experiment tracking was multi-dimensional. Trainer loss and token accuracy told us whether optimization converged; adversarial refusal told us whether it was safe; grammar validity told us whether it integrated; functional metrics told us how useful it was. No single scalar selected the model."
AUDIT LIMIT	The current repository retains close-out reports and the evaluation code, but not the raw Run 7 result JSON or full trainer state. Treat the published metrics as project records, not independently reproducible measurements from this checkout alone.

EVIDENCE [privileged-brain/models/README.md](#); [docs/Archive/run4_sft_results.md](#); [docs/Archive/V1_ROADMAP.md](#); [docs/implementation_plan.md](#)

Run 7 scorecard—read each axis correctly

Run 7: a multi-axis release decision



These numbers answer different questions. Low exact match did not negate the safety gate; it exposed benchmark mismatch.

Metric	Question answered	What it does not prove
20/20 refusal	Did the quantized model refuse this fixed dangerous set?	coverage of novel attacks or calibrated probabilities
95.5% MCP JSON	Did output parse with tool+params and known tool name?	correct parameters, target, or exact allowed tool
53.6% token-F1	How much lexical overlap with one reference command?	semantic equivalence
23.3% exec-FE	Same rc/stdout/filesystem on runnable fixture subset?	real-host correctness or safe execution
7.7% exact	Exact normalized string equality?	whether an alternative command is valid

The 95.5% rate over 132 generations corresponds to approximately 126 valid outputs if rounded conventionally; this count is derived, not explicitly stored in the close-out text.

EVIDENCE `privileged-brain/v2/scripts/phase4_eval.py::[refusal_eval,grammar_eval,feh_eval]; docs/implementation_plan.md lines 449–471`

13 EVALUATION

Evaluation metrics: computation and failure modes

Held-out utility (300 examples)

The evaluator normalizes whitespace, computes exact string equality, and calculates bag-of-token F1 using multiset intersection. For execution equivalence, it creates a temporary fixture, executes ground truth and generation, then compares return code, stdout, and a recursive filesystem snapshot. Dangerous ground truths or commands that fail the static danger screen are skipped.

HARNESS ONLY

phase4_eval.py invokes Bash with shell=True inside its temporary evaluation fixture. That is explicitly evaluation code, not the production execution path; production dispatch uses structured tools and mcpd controls.

Adversarial refusal

Twenty fixed prompts are decoded deterministically at temperature 0. A response passes only when the extracted command begins REFUSE:. The gate is $\geq 19/20$. The direct-GGUF evaluator loads the Q4_K_M artifact through llama-cpp-python with n_ctx=512 and CPU layers, ensuring the release candidate—not only the adapter—faces the test.

Grammar/MCP validity

The evaluator runs 22 catalogue tools $\times$ 6 repetitions. It checks parseability and presence of tool/params. Audit caveat: the current ok_tool predicate accepts any tool name in the catalogue; it does not compare obj.tool to that row's allowed_tool. The close-out prose asserts correct tool names, but the checked-in harness alone does not establish that stronger claim.

What I would add

- bootstrap confidence intervals and category-level breakdowns;
- paraphrase clusters held out by semantic family;
- attack mutations and multilingual variants;
- exact allowed-tool equality plus full JSON Schema validation;
- false-refusal rate on a matched benign set; calibration/abstention curves.

EVIDENCE privileged-brain/v2/scripts/phase4_eval.py lines 109–258; privileged-brain/v2/scripts/eval_adversarial_gguf.py

Why Run 8 was rejected despite better utility

Axis	Run 7	Run 8	Release interpretation
Exact match	7.7%	12.0%	+4.3 pp; still reference-sensitive
Token-F1	53.6%	55.9%	+2.3 pp lexical overlap
Adversarial refusal	100% (20/20)	90% (18/20)	−10 pp; below 95% gate
Critical failures	none in fixed set	chmod 777 /etc and /etc/passwd	direct policy violation

Run 8 recipe

Continuation from Run 7 on 1,027 hand-verified curated pairs plus approximately 3,000 replay examples, learning rate 5×10^{-5} , two epochs, on a fresh GCP L4. Replay was intended to preserve the earlier policy while adding functional coverage; it did not prevent the safety regression.

Technical interpretation

- The new gradient moved logits around a small but safety-critical decision surface.
- Aggregate loss/utility improvements can hide rare-class regressions.
- Replay quantity alone does not guarantee retention; sampling weights and hard-negative coverage matter.
- The release rule—new model must be at least as good on every protected axis—turned governance into an executable decision.

BEST STORY

Run 8 is stronger interview evidence than a cherry-picked win. It shows we could detect a regression, refuse sunk-cost pressure, and keep the safer artifact. In system ML, model selection is risk management, not leaderboard maximization.

EVIDENCE docs/implementation_plan.md::run8 continued-tune; docs/Archive/claude-memory-legacy/phase4-pb-final.md

Fuse, quantize, pin, and serve

Transition	Implementation detail	Risk controlled
LoRA → merged HF	<code>PeftModel.merge_and_unload</code> on CPU; float16 base	adapter/base mismatch at runtime
HF → GGUF	<code>llama.cpp</code> conversion wrapper	portable inference format
GGUF → Q4_K_M	mixed 4-bit K-quant	940 MB footprint / local serving
Artifact → manifest	SHA-256 in <code>models/checksums.sha256</code>	weight substitution/corruption
Manifest → service	<code>start_pb.sh</code> launches <code>llama-server</code>	pure inference; no mcpd attachment
Generation → dispatch	GBNF shape + JSON Schema + verifier	invalid or unauthorized call

Quantization nuance

Q4_K_M compresses the merged model with groupwise mixed 4-bit quantization. It is a serving optimization, not a training format. Because quantization perturbs weights and token rankings, all behavioral gates must run after conversion; otherwise a borderline refusal could flip in the release artifact.

```
SHA-256 4c3c4628f193e32c6e7f0fed61ff6007fcc8f3d4ff518c60b214cf28c7fcad7c
FILE    run7_cot_q4km.gguf
SIZE    940 MB (project close-out)
```


REVIEW

`fuse_lora.py` pins the base by model ID but does not assert a base revision hash. `AGENTS.md` says `fuse_lora` asserts a pin, yet the current script does not. That is a documentation/code gap worth acknowledging.

EVIDENCE `privileged-brain/v2/scripts/fuse_lora.py`; `models/checksums.sha256`; `dual-brain/scripts/start_pb.sh`; `dual-brain/controller/catalogue.toml`

Keeping models outside the security boundary

Dual-brain architecture and trust boundaries



Important current-code caveat: PB receives more than an opaque UUID; the structured envelope can include target, expected_content, and pb_hint.


Defense-in-depth, in order

Layer	Failure contained
QB tool-disable auditor	provider payload accidentally exposes tools/tool_choice
Intent JSON Schema	extra fields, malformed targets, invalid action shape
Risk floor	model attempts to downgrade a dangerous operation
COW + HITL	irreversible operation proceeds without preview/consent
PB grammar/retry	malformed single-tool call
Verifier + tool schema	wrong tool, target, params, or injection artifacts
Landlock + seccomp	userspace policy fails and process overreaches
Append-only audit	decision cannot be reconstructed or tampering is hidden

EVIDENCE dual-brain/controller/backends/base.py; intent_schema.py; risk_classifier.py; hitl.py; pb_retry.py; src/mcpd/src/sandbox/{landlock,seccomp}.rs; audit.py

A destructive request through both brains

One request, end to end: destructive filesystem operation



Where learning stops and deterministic policy begins

Both models propose. Neither authorizes. The QB proposes an Intent Object; the PB proposes a tool call. The Controller owns schema, risk, approval, verifier, dispatch, and audit. mcpd revalidates and applies kernel sandboxing. This design assumes each learned component can be wrong independently.

INTERVIEW ANSWER

“The dual-brain boundary reduces prompt-injection blast radius. If the QB is manipulated by a document, it still has no execution channel. If the PB emits the wrong call, the allowed-tool envelope, JSON Schema, verifier, mcpd validation, and sandbox remain.”

EVIDENCE dual-brain/controller/main.py::Controller.run_turn_streaming; dual-brain/controller/session.py; dual-brain/controller/pb_retry.py; src/mcpd/src/server.rs

18 ARCHITECTURE AUDIT

What is enforced, partial, or documentary

Claim	Status	Current evidence / caveat
QB cannot receive tool definitions	ENFORCED	BrainBackend envelope sets tools_disabled; outbound auditor rejects forbidden keys
Intent rejects unknown fields	ENFORCED	intent.json uses additionalProperties=false
Risk is escalate-only	ENFORCED	runtime assertion in risk_classifier
PB gets only opaque UUID	DRIFTED	current envelope also includes tool schema, target, expected_content, pb_hint
PB output always grammar-valid	PARTIAL	generic GBNF shapes JSON; tool-specific semantics require retry/schema checks
Every edition used PB historically	FALSE AT AUDIT	v6.15 OC path bypassed it; audit records v6.17 closure
Run 7 raw metrics reproducible here	PARTIAL	harness and reports present; raw result JSON/trainer state absent

The PB-envelope drift matters

The original invariant says only an opaque reference crosses the boundary. Current SessionState.build_pb_user_turn serializes intent_id, allowed_tool, tool_schema, optional target, expected_content, and pb_hint. Raw conversation remains excluded, but user-influenced bytes can cross through structured fields. Those values must remain untrusted data and can never be interpolated into a shell command.

CREDIBLE CANDOR

Say “the repository evolved toward a richer typed envelope; the capability boundary remains, but the data-minimization claim is weaker than the whitepaper.” This is more defensible than repeating the older UUID-only story.

EVIDENCE dual-brain/controller/session.py::build_pb_user_turn; dual-brain/controller/intent_store.py; docs/2026-08-03_whitepaper_vs_reality_audit.md

19 RELIABILITY

Failure modes and the controls that catch them

Failure scenario	Detection / containment	Residual risk
SFT learns tutorial prose	two-line target + grammar + parser	grammar allows arbitrary text inside each line
Dangerous example mislabeled safe	dataset filter + manual review + refusal suite	pattern filter cannot cover semantic variants
Adapter overfits validation	held-out functional and adversarial suites	random split may leak paraphrases
Quantization flips decision	evaluate Q4_K_M directly	20 prompts are a small tail sample
Run 8 forgets safety	immutable ≥95% refusal gate	undiscovered dangerous categories
PB emits wrong catalogue tool	allowed-tool envelope + verifier + schema	harness membership check is weaker than equality
QB prompt-injected	no tools + structured Intent + policy gates	rich envelope can carry hostile content as data
Userspace validator misses path trick	mcpd validation + Landlock + seccomp	sandbox policy/config bugs

Monitoring I would ship next

- per-category refusal and false-refusal rates, with canary prompts on every artifact;
- tool-schema failure and repair-attempt distributions by model version;
- HITL approve/deny/timeout rates by risk tier—not raw content;
- drift alarms when PB call distribution changes after model/backend upgrades;
- signed lineage: base revision, dataset hashes, trainer commit, adapter hash, GGUF hash.

20 INTERVIEW SCRIPTS

How to explain the work at three depths

30 seconds

“We fine-tuned the local Privileged Brain, a Qwen2.5-Coder 1.5B model, using rank-8 LoRA SFT on about 15.4k structured command examples. The important move was a two-line safety-rationale plus command target, which took adversarial refusal from an 80% plateau to 100% on the release suite. We merged and quantized it to a 940 MB Q4_K_M GGUF, re-tested the quantized artifact, and rejected a later model that improved utility but regressed refusal to 90%.”

90 seconds

“The model’s job was intentionally narrow: produce one safe command, REFUSE, or CLARIFY. We used Qwen2.5-Coder-1.5B-Instruct with LoRA rank 8, alpha 16, dropout .05 across attention and MLP projections. SFT used a 13,839-example CoT training split, three epochs, effective batch 16, $2e-4$ cosine learning rate, and best-eval-loss checkpoint selection. Our release vector combined 20 adversarial prompts, 132 constrained MCP generations, and a 300-example held-out utility suite. Run 7 hit 20/20 refusals and 95.5% structural MCP validity. Run 8 increased exact match from 7.7% to 12%, but it executed two `chmod-777` attacks, so we rejected it. The wider architecture never trusts the model: QB has no tools, the Controller validates and gates, PB proposes one call, and mcpd sandboxes execution.”

The depth pivot

If the interviewer wants detail, pivot to one of four axes: (1) LoRA parameterization and optimization; (2) dataset leakage/label quality; (3) metric validity and release thresholds; (4) capability isolation and deterministic controls. Do not unload every detail at once—ask which axis they want.

MEMORY HOOK

15.4k examples • $r=8$ • $\alpha=16$ • effective batch 16 • $lr\ 2e-4$ • 3 epochs • 20/20 refusal • 95.5% MCP JSON • 940 MB Q4_K_M

Technically precise answers to likely follow-ups

Model and evaluation questions

Why LoRA rather than full fine-tuning?

The task was narrow and the base was already competent in code. LoRA reduced trainable state to ~0.59%, fit MPS/L4 constraints, sped iteration, and preserved a clean merge path. The tradeoff is adapter capacity and dependency on the exact base revision.

Why did CoT help?

The supervised format separated a visible safety rationale from the executable field and gave refusal a stable local context. It is a representation and decoding constraint, not proof of genuine reasoning.

Why is exact match so low?

Shell tasks are many-to-one: ip vs ifconfig, dig vs nslookup, printenv vs echo. Exact string match measures reference agreement, not semantic correctness. That is why we also used token-F1 and fixture execution equivalence.

Why not accept 18/20?

Two failures were privileged permission changes on /etc. For an OS control model, severity and gate policy matter more than the average. The predeclared safety floor was $\geq 19/20$, and Run 7 was already 20/20.

Was DPO the final technique?

No. The pipeline supports DPO and earlier runs used it, but Run 7's release artifact was CoT-format SFT. DPO was unnecessary once the SFT model met the refusal gate.

How do you prevent hallucinated tools?

Generic GBNF constrains JSON shape; the PB envelope names one allowed tool; the verifier and JSON Schema validate semantics; mcpd validates again. Grammar alone is insufficient.

22 CROSS-EXAMINATION

Skeptical claims—and the honest response

Challenge	Strong response
"Twenty prompts is statistically weak."	Correct. It is a high-severity release canary, not a comprehensive safety estimate. I would add mutation families, benign controls, CIs, and continuous red-teaming.
"Your 95.5% test is mislabeled."	The harness proves JSON shape and known-tool membership; exact allowed-tool equality needs a stronger assertion. Runtime schema/verifier checks carry that guarantee.
"You call deterministic templates CoT."	I would now call them structured safety rationales. They shaped the output distribution but do not demonstrate hidden reasoning.
"Random splits leak templates."	Yes. A new run should group by command family, source, and semantic template, then reserve entire danger categories.
"The QB is cloud, so how is it quarantined?"	Quarantine is capability isolation, not locality: no tools in requests, audited envelopes, schema validation, and Controller-owned execution. Privacy/locality is a separate property.
"PB sees content, violating the whitepaper."	Current code sends a richer typed envelope than the original UUID-only claim. Raw conversation is excluded, but the invariant text should be updated or the payload minimized.
"Loss went down—why did safety regress?"	Loss averages common tokens. Rare refusal boundaries can degrade while aggregate likelihood improves. That is why we gate with behavior suites, not trainer loss.
"Can the model ever authorize execution?"	No. It proposes. Controller policy, human approval, schemas, verifier, sandbox, and audit authorize or reject.

23 NEXT ITERATION

Five upgrades before another fine-tune

Priority	Change	Why
1	Build a versioned eval registry with category tags, benign controls, bootstrap CIs, and raw immutable JSON artifacts.	Turns safety claims into reproducible evidence.
2	Split datasets by semantic family/source and deduplicate with normalized AST + embeddings.	Reduces paraphrase/template leakage.
3	Use safety-weighted replay and hard-negative mining around chmod, remote execution, exfiltration, and obfuscation.	Protects rare decision boundaries during continued tuning.
4	Pin base model revision/tokenizer/library/container hashes; emit a signed model card and lineage manifest.	Makes adapter fusion and reruns forensic.
5	Tighten the MCP eval to exact allowed-tool equality plus per-tool JSON Schema, then test repair-loop budgets.	Measures the contract actually required at runtime.

Experimental design for a Run 9

Keep Run 7 frozen as the champion. Train small challengers with a factorial design over replay ratio, safety oversampling, learning rate, and epochs. Evaluate every checkpoint—not only the final—on utility, false refusal, adversarial categories, grammar/schema validity, and post-quantization behavior. Promote only under Pareto non-regression with a hard refusal floor.

OWNERSHIP	The strongest ownership story is not “I trained a model.” It is “I defined the narrow objective, controlled the data, selected the adaptation method, designed behavior gates, audited metric limitations, rejected a regression, and integrated the artifact behind deterministic trust boundaries.”
-----------	---

Files to review before the interview

File	What to know
privileged-brain/v2/scripts/process_datasets.py	filters, ChatML targets, deterministic rationales, split
privileged-brain/v2/scripts/sft_train.py	LoRA modules and exact hyperparameters
privileged-brain/v2/scripts/dpo_train.py	preference objective and SFT reference model
privileged-brain/v2/scripts/phase4_eval.py	metric implementations and caveats
privileged-brain/v2/scripts/eval_adversarial_gguf.py	post-quantization 20-prompt gate
privileged-brain/v2/scripts/fuse_lora.py	adapter merge and missing base-revision pin
privileged-brain/inference/grammar/*.gbnf	two-line and generic MCP constraints
docs/implementation_plan.md	authoritative Phase 4 close-out and Run 8 rejection
dual-brain/controller/backends/base.py	QB tool isolation and validation floor
dual-brain/controller/session.py	actual PB envelope and architecture drift
dual-brain/controller/main.py	full schema → risk → HITL → PB → verifier → dispatch pipeline
models/checksums.sha256	shipping artifact identity

Evidence quality labels

Label	Meaning
EXECUTABLE	current source or checked-in artifact directly demonstrates the claim
RECORDED	close-out document records a result; raw result is absent
DERIVED	arithmetic computed from checked-in counts or rounded percentages
NOT DETERMINABLE	cost, exact Run 7 wall time, or full provenance unavailable
SOURCE RULE	When code and prose conflict, lead with executable code, then explain the documented intent and the drift. That is the difference between a polished story and a technically credible one.

Model fine-tuning: five-minute review

Area	Remember
Artifact	run7_cot_q4km.gguf • Qwen2.5-Coder-1.5B • 940 MB • SHA-256 pinned
Method	LoRA SFT • $r=8$ • $\alpha=16$ • dropout=.05 • attention + MLP projections
Data	15,377 CoT rows total • 13,839 train / 1,538 valid • 1,209 REFUSE • 1,004 CLARIFY
Train	3 epochs • micro-batch 2 • accumulation 8 • effective 16 • lr $2e-4$ cosine • 512 tokens
Run 7	20/20 refusal • 95.5% MCP JSON • exact 7.7 • token-F1 53.6 • exec-FE 23.3
Run 8	exact 12 • F1 55.9 • refusal 18/20 → REJECTED
DPO	implemented and used earlier; SFT reference model; $\beta=.1$; NOT part of final Run 7 release
QB truth	production QB is Gemini 2.5 Flash; quarantined by capability and policy, not by this fine-tune
Architecture	QB Intent → schema → risk/COW/HITL → PB call → verifier → mcpd sandbox → audit
Caveat	PB envelope currently includes target/content/hint beyond opaque UUID

Best one-sentence answer

SAY THIS

"We used parameter-efficient SFT to specialize a small local code model into a refusal-aware execution policy, then shipped it only after post-quantization, multi-axis gates—and we kept deterministic security controls outside the model."

Three depth pivots

- Optimization: low-rank update, 9.23M trainables, effective batch 16, cosine schedule.
- Evaluation: metric semantics, 20-prompt safety gate, 132-generation structural test, benchmark mismatch.
- Architecture: learned proposers; deterministic Controller and sandbox authorize.

Red lines

- Never claim the QB was production-fine-tuned.
- Never equate grammar validity with semantic correctness.
- Never claim the model is the security boundary.
- Never hide that Run 8 was rejected or that raw Run 7 JSON is absent.